
Missing the Point: Non-Convergence in Iterative Imputation Algorithms

Sociological Methods & Research
2020, Vol. 49(?):1–19
© The Author(s) 2020
Article reuse guidelines:
sagepub.com/journals-permissions
DOI: 10.1177/ToBeAssigned
journals.sagepub.com/home/smr



H. I. Oberman¹

Abstract

(**Rewrite:**) Multiple imputation by chained equations (MICE) is a widely used tool to accommodate missing data. While empirical evidence supports the validity of inferences obtained using mice, there is no consensus on the convergence properties of the method. This paper provides insight into non-convergence of mice algorithms.

Keywords

MICE; convergence

Introduction

Anyone who analyzes person-data may run into a missing data problem. Missing data is not only ubiquitous but also tedious. If a dataset contains just one incomplete observation, statistical inferences are undefined and will not produce any results. To circumvent this, many statistical packages employ list-wise deletion by default (i.e., ignoring incomplete observations). Unfortunately, this *ad hoc* solution may yield wildly invalid results (Van Buuren 2018). An alternative is to ‘impute’ (i.e., fill in) the missing values in the incomplete observations. Subsequently, statistical inferences can be performed on the completed dataset. By repeating this process several times, a distribution of plausible results may be obtained, which reflects the uncertainty in the data due to missingness. This technique is known as ‘multiple imputation’ (MI; Rubin 1976). MI yields unbiased and confidence valid estimates of the true—but missing—data inference under many circumstances (Van Buuren 2018).

¹Department of Methodology and Statistics, Utrecht University, Utrecht, The Netherlands

Corresponding author:

Hanne Oberman, Sjoerd Groenman building, Utrecht Science Park, Utrecht, The Netherlands.

Email: h.i.oberman@uu.nl

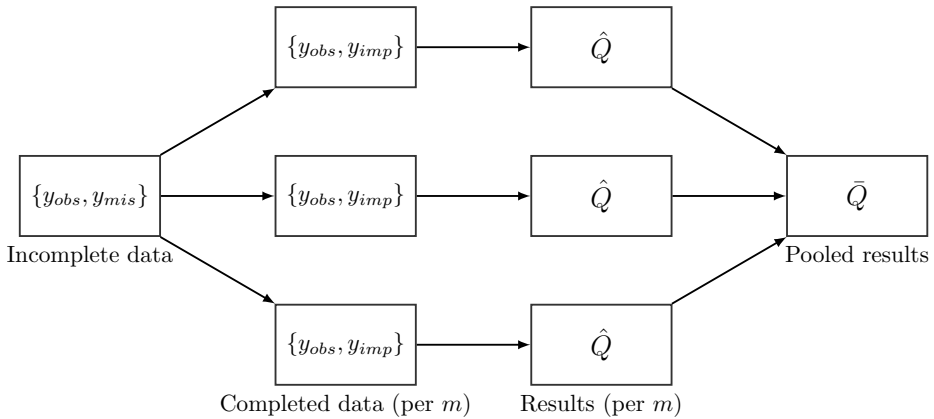


Figure 1. Scheme of the main steps in multiple imputation.

Move this after terminology?? Figure 1 provides an overview of the steps involved with MI—from incomplete data, to M completed datasets, to M estimated quantities of interest (\hat{Q} s), to a single pooled estimate \bar{Q} . Missing data in y is ‘imputed’ (i.e., filled in) M times. The imputed data (y_{imp}) is combined with the observed data (y_{obs}) to create M completed datasets. On each completed dataset, the analysis of scientific interest (the ‘substantive model’, or ‘complete data model’) is performed. The quantity of scientific interest (e.g., a regression coefficient) is denoted with Q . Since Q is estimated on each completed dataset, M separate \hat{Q} -values are obtained. These M values are combined into a single pooled estimate \bar{Q} . **ALSO: make clear that iteration happens within the first step, between incomplete and completed data.**

A popular method to obtain imputations is to use the ‘Multiple Imputation by Chained Equations’ algorithm, shorthand ‘MICE’ (Van Buuren and Groothuis-Oudshoorn 2011). MICE is an iterative algorithmic procedure to draw imputations from the posterior predictive distribution of the missing values. This introduces a potential threat to the validity of the imputations: What if the algorithm has not converged? Are the implications then to be trusted? And can we rely on the inference obtained on the completed data? These are all open questions, because the convergence properties of iterative imputation algorithms have not been systematically studied (Van Buuren 2018). Moreover, there is no scientific consensus on the convergence properties of MI algorithms (Takahashi 2017). Some default MICE techniques (e.g., ‘predictive mean modeling’) might not yield converged states at all (Murray 2018). Therefore, algorithmic convergence should be monitored carefully.

Currently, the recommended practice for evaluating convergence is to visually inspect imputations for signs of non-convergence. This method is insufficient on two counts: 1) it may be challenging to the untrained eye, and 2) only severely pathological cases of non-convergence may be diagnosed (Van Buuren 2018, § 6.5.2). Therefore, a quantitative, diagnostic evaluation of convergence would be preferred.

Monitoring convergence of iterative imputation algorithms diagnostically is challenging. Iterative imputation algorithms such as MICE are Markov chain Monte Carlo (MCMC) methods. In MCMC methods, convergence is not from a scalar to a point but from one distribution to another. The values generated by the algorithm (e.g., imputed values) will vary even after convergence. Therefore, the aim of convergence diagnostics for MCMC methods is not to establish the point at which convergence is reached, but to monitor signs of non-convergence. Several of such diagnostics exist for MCMC methods, but it is not known whether these are appropriate for MICE.

In this paper, we investigate how non-convergence in iterative imputation algorithms may be diagnosed, how well these methods perform, and at which point convergence may safely be assumed. For reasons of brevity, we only focus on the MI algorithm implemented in the world-leading MI software: the R (R Core Team 2020) package `mice` (Van Buuren and Groothuis-Oudshoorn 2011). The convergence properties of this MI algorithm are investigated through model-based simulation. The results of this simulation study are guidelines for assessing convergence of MI algorithms, which will aid applied researchers in drawing valid inference from incomplete datasets.

Include sub-questions?: How can non-convergence be diagnosed? Are common MCMC non-convergence diagnostics appropriate for MICE? And if so, which threshold should be used to diagnose non-convergence? How many iterations are sufficient/needed to be able to diagnose non-convergence? Are the default number of iterations sufficient (i.e., 5 in `mice`, 10 in SPSS and Stata, 30 in `mi`)?

Terminology

Let y denote an $n \times p$ matrix containing the data values on p variables for all n units in a sample. The data value of unit i ($i = 1, 2, \dots, n$) on variable j ($j = 1, 2, \dots, p$) may be either observed or missing. The collection of observed data values in y is denoted by y_{obs} ; the missing part of y is referred to as y_{mis} . For each datapoint in y_{mis} , we sample $M \times T$ times plausible values, where M is the number of imputations ($m = 1, 2, \dots, M$) and T is the number of iterations ($t = 1, 2, \dots, T$). The collection of samples between the initial value (at $t = 1$) and the final imputed value (at $t = T$) will be referred to as an ‘imputation chain’. **Add: missingness percentage = proportion of cases with one or more missing values, Q is quantity of scientific interest:** “A scientific estimand Q is a quantity of scientific interest that we can calculate if we would observe the entire population” (Van Buuren 2018, par 2.3.1), and θ s are scalar summaries of interest in the iterative algorithm (e.g., chain means; the average of the imputed values in each imputation chain).

How may non-convergence be diagnosed?

The two challenges for convergence of iterative algorithms are mixing and stationarity (Gelman et al. 2013). Without mixing, imputation chains do not intermingle nicely. Without stationarity, there is trending within imputation chains, see Figure 2. **Explain what we see, namely example by Van Buuren (2018) reproduced, showing the traceplots of chain means for some variable. The first plot is typical convergence**

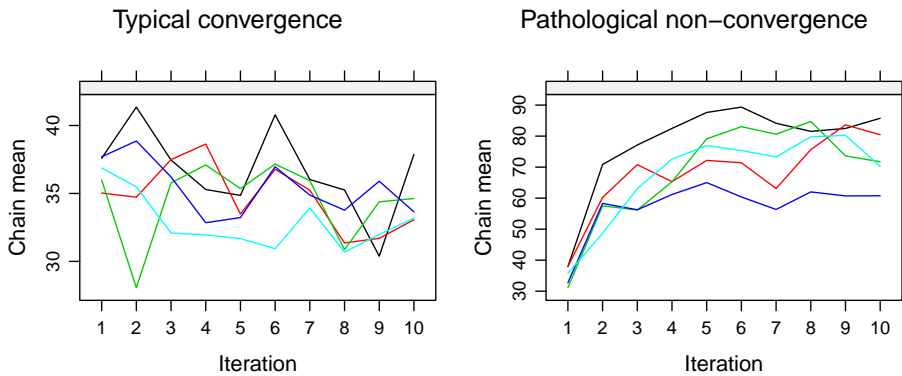


Figure 2. Typical convergence versus pathological non-convergence.

of MICE, the second is pathological non-convergence because of a mis-specified imputation model. Each line is an imputation. In the first plot, the lines intermingle nicely and there is little trending. In the second, there is a lot of trending and the lines do not cross much. Importantly, the chain means at the last iteration (the imputed value per m) are very different between the two plots. The algorithm with the mis-specified model yields imputed values that are on average a magnitude two larger than those of the typically converged algorithm. This shows the importance of reaching converged states in iterative imputation algorithms.

Non-stationarity within chains may be diagnosed with e.g., autocorrelation (AC; Schafer 1997; Gelman et al. 2013), numeric standard error ('MC error'; Geweke 1992), or Raftery and Lewis's (1991) procedure to determine the effect of trending on the precision of estimates. A widely used diagnostic to monitor mixing between chains is the potential scale reduction factor \hat{R} ('Gelman-Rubin statistic'; Gelman and Rubin 1992). With a recently proposed adaptation, \hat{R} might also serve to diagnose non-stationarity, but this has not yet been thoroughly investigated (Vehtari et al. 2019). Therefore, mixing and stationarity will be evaluated separately in this study. As recommended (e.g., Cowles and Carlin 1996, p. 898), AC and \hat{R} will be used.

Potential scale reduction factor

To define \hat{R} , we follow notation by (Vehtari et al. 2019, p. 5). Let M be the total number of chains, T the number of iterations per chain, and θ the scalar summary of interest (e.g., chain mean or chain variance). For each chain ($m = 1, 2, \dots, M$), we estimate the variance of θ , and average these to obtain within-chain variance W .

$$W = \frac{1}{M} \sum_{m=1}^M s_m^2, \text{ where } s_m^2 = \frac{1}{T-1} \sum_{t=1}^T \left(\theta^{(tm)} - \bar{\theta}^{(\cdot m)} \right)^2.$$

We then estimate between-chain variance B as the variance of the collection of average θ per chain.

$$B = \frac{T}{M-1} \sum_{m=1}^M \left(\bar{\theta}^{(\cdot m)} - \bar{\theta}^{(\cdot \cdot)} \right)^2, \text{ where } \bar{\theta}^{(\cdot m)} = \frac{1}{T} \sum_{t=1}^T \theta^{(tm)}, \bar{\theta}^{(\cdot \cdot)} = \frac{1}{M} \sum_{m=1}^M \bar{\theta}^{(\cdot m)}.$$

From the between- and within-chain variances we compute a weighted average, $\widehat{\text{var}}^+$, which over-estimates the total variance of θ **remove or explain why**. \widehat{R} is then obtained as a ratio between the over-estimated total variance and the within-chain variance:

$$\widehat{R} = \sqrt{\frac{\widehat{\text{var}}^+(\theta|y)}{W}}, \text{ where } \widehat{\text{var}}^+(\theta|y) = \frac{N-1}{N}W + \frac{1}{N}B.$$

We can interpret \widehat{R} as potential scale reduction factor since it indicates by how much the variance of θ could be shrunk down if an infinite number of iterations per chain would be run (Gelman and Rubin 1992). This interpretation assumes that chains are ‘over-dispersed’ at $t = 1$, and reach convergence as $T \rightarrow \infty$. Over-dispersion implies that the initial values of the chains are ‘far away’ from the target distribution and each other. When all chains sample independent of their initial values, the mixing component of convergence is satisfied, and \widehat{R} -values will be close to one. High \widehat{R} -values thus indicate non-convergence. The conventionally acceptable threshold for convergence was $\widehat{R} < 1.2$ (Gelman and Rubin 1992). More recently, Vehtari et al. (2019) proposed a more stringent threshold of $\widehat{R} < 1.01$.

Autocorrelation

Following the same notation, we define autocorrelation as the correlation between two subsequent θ -values within the same chain (Lynch 2007, p. 147). In this study, we only consider AC at lag 1, i.e., the correlation between the t^{th} and $t + 1^{th}$ iteration of the same chain.

$$AC = \left(\frac{T}{T-1} \right) \frac{\sum_{t=1}^{T-1} (\theta_t - \bar{\theta}^{(\cdot m)})(\theta_{t+1} - \bar{\theta}^{(\cdot m)})}{\sum_{t=1}^T (\theta_t - \bar{\theta}^{(\cdot m)})^2}.$$

We can interpret AC -values as a measure of stationarity. If AC -values are close to zero, there is no dependence between subsequent samples within imputation chains. Negative AC -values indicate divergence within imputation chains. **Subsequent sampled values within each imputation chain are less alike**. Positive AC -values indicate recurrence. If θ -values of subsequent iterations are similar, trending may occur. Negative AC -values show no threat to the stationarity component of convergence. On the contrary even—negative AC -values indicate that θ -values of subsequent iterations diverge from one another, which may increase the variance of θ and speed up convergence. As convergence diagnostic, the interest is therefore in positive AC -values. **Maybe remove:** Moreover, the magnitude of AC -values may be evaluated statistically, but that is outside of this note’s scope.

In practice/Answer

To assess whether \hat{R} and AC may be appropriate non-convergence diagnostics for MICE algorithms, the following condition must hold when applied on the two algorithms plotted above: the methods indicate worse performance for the mis-specified model with pathological non-convergence (i.e., higher \hat{R} - and AC -values than the typical performance). And additionally the methods should reflect the increasing convergence in the typical convergence situation as the number of iterations goes up.

In figure 3A (add panel labels), the chain means are plotted again, now together. So are the non-convergence diagnostics. Panel B shows \hat{R} as computed by implementing [Vehtari et al. \(2019\)](#)'s recommendations. As required, \hat{R} indicates less signs of non-convergence as the number of iterations goes up in the typical convergence situation. The superior performance of the typical convergence over the pathological non-convergence is less prominent, even flipped for $t < 4$. Panel C displays the AC as computed with the R function `stats::acf()`. When we look at this panel, we must conclude that something weird is going on. The AC -values indicate equal performance (up-to $t = 5$) for the typical convergence and the pathological non-convergence, while there is obvious trending in the latter. Moreover, the best convergence (as indicated by the lowest AC -value) is observed at $t = 2$, but looking at the chain means in panel A, there should be some signs of trending up-to iteration number seven. After consulting the documentation on `stats::acf()`, we conclude that this implementation of AC is not suitable for iterative imputation algorithms. There is a correction factor for a mathematical shortcut that works in the limit. According to... the function works when the number of iterations > 50 . The default number of iterations in `mice`, however, is 5. Therefore, we compute AC manually, see panel D. The AC values in this plot do meet the requirements (...) and will therefore be used in the simulation study.

Simulation Hypothesis

Convergence is reached when there is no dependency between subsequent iterations of imputation chains ($AC = 0$), and chains intermingle such that the only difference between the chains is caused by the randomness induced by the algorithm ($\hat{R} = 1$). We expect that a completely converged state will not be reached.

This study evaluates whether \hat{R} and AC could diagnose convergence of multiple imputation algorithms. We assess the performance of the two convergence diagnostics against the recommended evaluation criteria for MI methods (i.e., average bias, average confidence interval width, and empirical coverage rate across simulations; [Van Buuren 2018](#), § 2.5.2). That is, there is no baseline measure available to evaluate performance against.

Based on an empirical finding ([Lacerda et al. 2007](#)), we hypothesize that \hat{R} will over-estimate non-convergence of MI algorithms **explain why?**. The threshold of $\hat{R} < 1.01$ will then be too stringent for diagnosing convergence. This over-estimation may, however, be diminished because \hat{R} can falsely diagnose convergence if initial values of

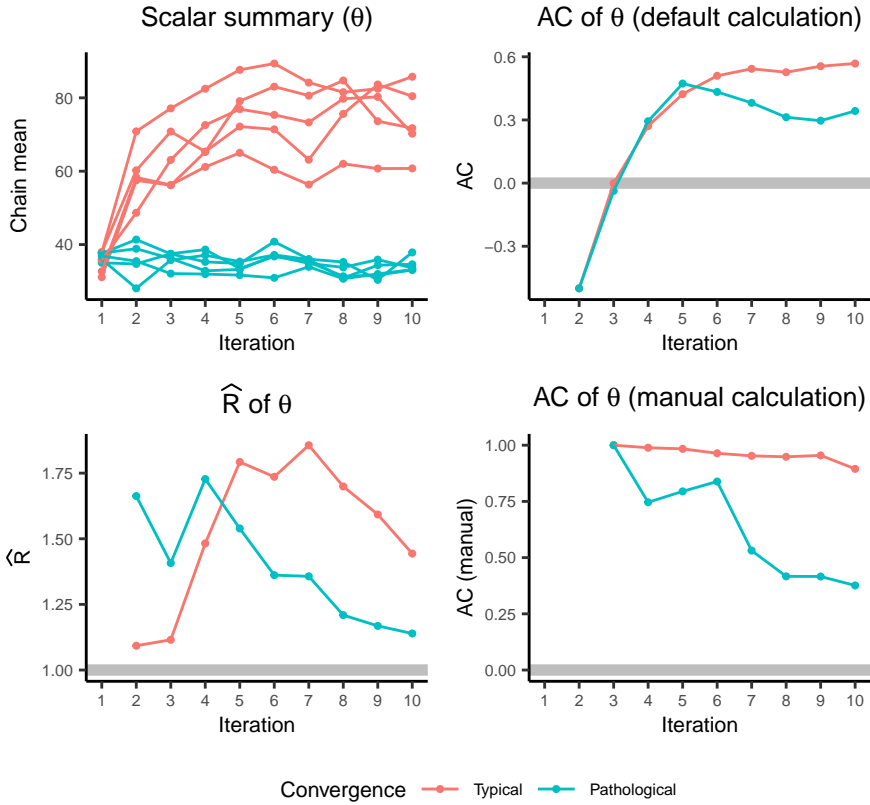


Figure 3. Convergence diagnostics pathological non-convergence versus typical convergence.

the algorithm are not appropriately over-dispersed (Brooks and Gelman 1998, p. 437). In mice, initial values are chosen randomly from the observed data. Therefore, we cannot be certain that the initial values are over-dispersed. We expect this to have little effect on the hypothesized performance of \hat{R} explain that the randomness induced by the MICE algorithm would take care of this??. (Add actual hypothesis: we thought that mice would converge sooner than \hat{R} would indicate that it did) No hypothesis was formulated about the performance of AC as convergence diagnostic. Added later: High AC-values are implausible in MI procedures. That is, the randomness induced by the MI algorithm effectively mitigates the risk of dependency within chains. But this doesn't seem the case in practice??

Simulation study

How to diagnose non-convergence?

The two challenges of convergence are stationarity and mixing. Mixing and stationarity can be inspected visually, by evaluating traceplots of scalar summaries of interest (θ s; e.g., chain means and chain variances). As Van Buuren (2018) describes, users can also specify a model-specific scalar summary (e.g., a regression coefficient). A user-specified scalar summary, however, is not universal to all complete data problems. Therefore, as inspired by (MacKay and Mac Kay 2003), we will propose a θ that summarizes the multivariate state space of the algorithm. Namely, the first eigenvalue of the variance-covariance matrix of the M completed datasets.

The methods under evaluation are \hat{R} and AC . These will be computed for the three scalar summaries of interest: chain means (i.e., per imputation, per variable: on $y_{imp,m,j}$), chain variances (i.e., per imputation, per variable: on $y_{imp,m,j}$), and the first eigenvalue of the variance-covariance matrix across imputations (i.e., per imputation: on $\{y_{obs}, y_{imp,m}\}$) **add user-specified θ : estimated regression coefficient per imputation.**

The performance of two methods (\hat{R} and AC) on each of these scalar summaries of interest will be evaluated with several estimands (quantities of scientific interest; Q s). It is assumed that when \bar{Q} (the pooled result across imputations) is an unbiased (and confidence valid) estimate of Q , the algorithm is sufficiently converged. For each scalar θ , one estimand Q is defined that may be of interest in empirical research. The \hat{R} and AC values with chain means as θ s are evaluated against the bias in univariate mean estimates (**all variables or just for Y??**). Similarly, the performance measure for \hat{R} and AC applied to chain variances is the bias in estimated standard deviations. To evaluate the performance of \hat{R} and AC on the eigenvalues, we will use bias in the coefficient of determination, R^2 . Additionally, we will evaluate the estimated regression coefficients, the coverage rate of the CI95% of the regression estimates, and the CI length (**see definitions in old Methods**). **remove the term performance measure here and add an extra paragraph to define performance measures as bias in all estimands, and coverage rate and CI length of regression coefficients.**

Move this to Results: As expected, conditions with a higher proportion of missingness and/or a lower number of iterations show more signs of non-convergence, as indicated by more extreme bias in the estimated Q s. Roughly speaking this means that MICE is indeed not converged as $t = 1$, and converges gradually as t increases. The point at which an additional iteration does not lead to an improvement of the estimates depends on the difficulty of the missingness problem—between XYZ and XYZ. This means that the algorithm has converged sufficiently (**under the current specifications**).

As required for a method to diagnose non-convergence, the values of \hat{R} follow a trend similar to the performance measures. \hat{R} values are generally lower in conditions with a higher number of iterations, and somewhat higher in conditions with a higher percentage of missing data. **Decide where to specify the two calculations of AC:** Initially, autocorrelation did not show a decrease with an increasing number of iterations.

But, after replacing the ACF function with a manual calculation for autocorrelation, the autocorrelation values were indeed decreasing with a higher number of iterations.

Evaluation with the performance measures shows that \hat{R} and autocorrelation are conservative: they indicate signs of non-convergence in conditions where \bar{Q} s are unbiased and or confidence valid estimates of Q .

When to diagnose non-convergence?

Upon convergence, $\hat{R} = 1$ and $AC = 0$, but in MCMC algorithms these are very unlikely (because it's a distribution, not a point to converge to). In practice, thresholds are used to diagnose non-convergence. Usually, non-convergence is diagnosed when $\hat{R} > 1.2$ or 1.1 or even 1.01. And a t-test is performed to assess whether AC is significantly different from zero. Performance measures are the same as above: unbiased, confidence valid estimates.

Move this to Results: As expected, complete convergence ($\hat{R} = 1$ and $AC = 0$) is not observed. But what about the thresholds for practice? $\hat{R} < 1.2$ is not stringent enough, because conditions in which \hat{R} is smaller than 1.2 do not all yield unbiased estimates of the Q s. Moreover, there is a dip in the \hat{R} values, after which $\hat{R} > 1.2$ again. If the algorithms would be terminated at the iteration where $\hat{R} < 1.2$ occurs for the first time, the increase in \hat{R} values might be missed. The threshold $\hat{R} < 1.1$ is somewhat conservative in comparison with the performance measures. But it seems OK. $\hat{R} < 1.01$ is too stringent compared to the performance measures. And it is not obtained for any number of iterations (**but may be necessary in more complex complete data models???**).

Start old methods section

Add to this section: 1) explicit mention of simulation conditions; 2) emphasize that the plots are averages across repetitions, not within MICE; 3) sample effects due to single complete dataset Convergence of the MICE algorithm is investigated through model-based simulation in R (version 3.6.3; [R Core Team 2020](#)). The simulation set-up is summarized in the pseudo-code below. The complete R script of the simulation study is available from github.com/gerkovink/shinyMice.

```
# pseudo-code of simulation
1. simulate data
for (missingness proportions 5%, 25%, 50%, 75% and 95%)
  for (number of simulation runs from 1 to 1000)
    2. create missingness
    for (number of iterations from 1 to 100)
      3. impute missingness
      4. compute convergence diagnostics
      5. perform analysis of scientific interest
      6. pool results across imputations
      7. compute performance measures
```

8. aggregate outcomes across simulation runs
9. combine outcomes of all missingness proportions

The aim of the simulation study is to evaluate the impact of inducing non-convergence by: 1) terminating the MICE algorithm at different imputation chain lengths ($t = 1, 2, \dots, 100$), and 2) varying the missingness proportions ($p_{miss} = .05, .25, .50, .75, .95$). The assumption underlying the different number of iterations is that the algorithm generally does reach convergence at $t = 1$, because the initial values are sampled randomly from the set of observed datapoints. **As the number of iterations goes up, the imputation chains will become independent of the initial values until the point at which an extra added iteration does not lead to a more converged state.** The second set of experimental conditions under consideration—the missingness proportions—are chosen to reflect the difficulty of the missingness problem. The inherent assumption is that low missingness proportions will lead to quick algorithmic convergence, since there is a lot of information in the observed data. Higher missingness proportions then cause slower convergence. However, if the fraction of missing information is very high, there is so little information in the data that the random component in the algorithm will take the overhand and a stable but very uncertain (high variance) point will be reached.

The data-generating mechanism is a multivariate normal distribution, representing person data on three independent variables (from an unspecified social scientific field of study). Let

$$\begin{pmatrix} X_1 \\ X_2 \\ X_3 \end{pmatrix} \sim \mathcal{N} \left[\begin{pmatrix} 12 \\ 3 \\ 0.5 \end{pmatrix}, \begin{pmatrix} 4 & 4 & 1.8 & 0 \\ 4 & 16 & 4.8 & 0 \\ 1.8 & 4.8 & 9 & 0 \end{pmatrix} \right].$$

In this study, sampling variance is not of interest. Therefore, a single complete set may serve as comparative truth in all simulation runs (**Vink**). A finite population of $N = 1000$ is simulated using the function `mvtnorm::rmvnorm()`. Subsequently, a fourth variable is constructed to serve as dependent variable in a multiple linear regression problem. Let

$$Y_i = 1 + 2X_{1i} + .5X_{2i} - X_{3i} + \epsilon_i,$$

where $i = 1, 2, \dots, N$ and $\epsilon \sim \mathcal{N}(0, 100)$.

We consider four quantities of scientific interest ('conceptual estimands'; **Morris et al. 2019**), namely... solve a multiple linear regression problem, where dependent variable Y is regressed on independent variables X_1 , X_2 and X_3 (**check notation with betas here, and add what the quantity/-ies of scientific interest is/are! Namely, the descriptive statistics of all variables (mean and standard deviation), β_2 , and the percentage of variance in Y explained (coefficient of determination; R^2)**):

$$Y \sim \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3.$$

The complete data is ‘amputed’ once for each simulation repetition with function `mice::ampute()`. **(change this to reflect current simulation set-up with 5, 25, 50, 75, and 95% of cases having missing data:** The missingness is univariate, and the probability to be missing is the same for all four variables, namely 20% (`prop = 0.8`, `mech = "MCAR"`). This leaves 20% of the rows completely observed). This study only considers only an MCAR missingness mechanism. Therefore, results may not be extrapolated to other missing data problems. Proper performance of the convergence diagnostics under MCAR is necessary but not sufficient to demonstrate the appropriateness of \hat{R} and AC as convergence diagnostics. Convergence diagnostics should therefore at least apply to the MCAR situation, before more complex missingness should be explored.

Missing datapoints in y are imputed with the function `mice::mice()`. All MI procedures are performed with Bayesian linear regression imputation (`method = "norm"`), and five imputation chains (`m = 5`). The number of iterations varies between simulation conditions (`maxit = 1, 2, \dots, 100`). Each repetition of the simulation thus starts from 1 complete dataset, then splits into 5 missingness conditions ($p = .05, .25, .5, .75, .95$), each of the amputed datasets is imputed 5 times ($m = 1, 2, \dots, 5$) for each iteration condition ($t = 1, 2, \dots, 100$). From these $5 \times 5 \times 100$ imputations, we extract the θ s to apply \hat{R} and AC on (**i.e., chain means, chain variances, β s per imputation, and the first eigenvalue of the variance-covariance matrix per imputation**). We calculate \hat{R} by implementing Vehtari et al.’s (2019) recommendations, and AC as the correlation between the t^{th} and the $(t + 1)^{th}$ iteration. For each θ we estimate the corresponding scientific estimand Q .

The estimator for each Q is \bar{Q} —the pooled aggregate of the (\hat{Q})s across imputations. To calculate the \bar{Q} s, we combine the observed data (y_{obs}) and the imputed data for each imputation ($y_{imp,m}$) with the function `mice::complete()`. Descriptive statistics are computed as the average across imputations for each variable (μ_j and σ_j , where $j = Y, X_1, X_2, X_3$). Estimated regression coefficients are obtained with the function `stats::lm()` for each imputation, and then pooled conform Vink and van Buuren (2014). The coefficient of determination is estimated for each imputation, and pooled using `mice::pool.r.squared()`.

The performance of \hat{R} and AC is assessed by comparing \bar{Q} s with Q s. For each Q , we compute bias as $\bar{Q} - Q$. For $Q = \beta_j$, we also compute the empirical coverage rate (CR). CR is defined as the percentage of simulation repetitions in which the 95% confidence interval (CI) around \bar{Q} covers the true estimand Q . Let

$$CI = \bar{Q} \pm t_{(M-1)} \times SE_{\bar{Q}},$$

where $t_{(M-1)}$ is the quantile of a t -distribution with $M - 1$ degrees of freedom, and $SE_{\bar{Q}}$ is the square root of the pooled variance estimate. **Remove this?** Under-estimating the variance of \bar{Q} may yield spurious inferences. Confidence interval width (CIW) is defined as the difference between the lower and upper bound of the 95% confidence interval (CI) around \bar{Q} . From bias and CIW, we calculate empirical coverage rates.

Coverage rate is the proportion of simulations in which Q is between the bounds of the CI95% around \bar{Q} .

Old Results section

For reasons of brevity, we only discuss the convergence diagnostics for the Q with the worst performance in terms of bias. For the descriptive statistics, the magnitude of the bias was the largest in Y (i.e., $j = Y$ in μ_j and σ_j). For the $Q = \beta_j$, $j = 2$ showed the worst performance (i.e., the effect of X_1 on Y).

(Add more info about figure legends and axes.)

Univariate estimates and convergence diagnostics

The bias in the estimates of the variable means shows little to no difference between simulation conditions. It doesn't seem to matter how many iterations you use in the mice algorithm, the estimates are unbiased. Similarly, the bias in the estimated variances is more or less stable across simulation conditions. These univariate quantities appear to be unaffected by the number of iterations.

When applied to the imputation chain means, \hat{R} indicates that the mice algorithm does not reach a converged state ($\hat{R} = 1$) in any of the simulation conditions. Neither is the most recent recommended threshold reached ($\hat{R} < 1.01$). **This is only for 75% missingness:** The conventional \hat{R} threshold of 1.2 is reached in simulation conditions $T = 3$ and $T > 6$. (With the default number of iterations (maxit = 5), this dip in \hat{R} values would be spotted, so it is no problem.) The point at which an extra iteration does not seem to improve the \hat{R} value is around $T = 30$.

Autocorrelations indicate no sign of trending (**NOT correct anymore, fluke of the stats::acf() function!!**) within imputation chains. In most simulation conditions AC is smaller than or about equal to zero. Across simulation conditions, however, the autocorrelation curve does not trend towards 0. Autocorrelation values plateau off at a value of around .1. This is a small positive autocorrelation, which would indicate some trending within chains.

The \hat{R} and AC values for the imputation chain variances show equal trends to the chain means and are therefore not discussed separately. Taken together, univariate estimates seem robust across simulation conditions. There is no clear effect of the number of iterations on the bias in these estimates, while the convergence diagnostics indicate that the algorithm did not reach a completely converged state (yet).

Multivariate estimates

This is only true for 75% missingness: There is a clear bias in the regression estimates in simulation conditions where the number of iterations is smaller than four. In simulation conditions where $T > 5$ there is little to no bias in the estimated regression coefficients. In most simulation conditions nominal coverage is obtained (i.e., coverage rates of .95) for the confidence intervals around the regression coefficients. Conditions with only one or two iterations show some under-coverage. Since confidence interval width is stable across

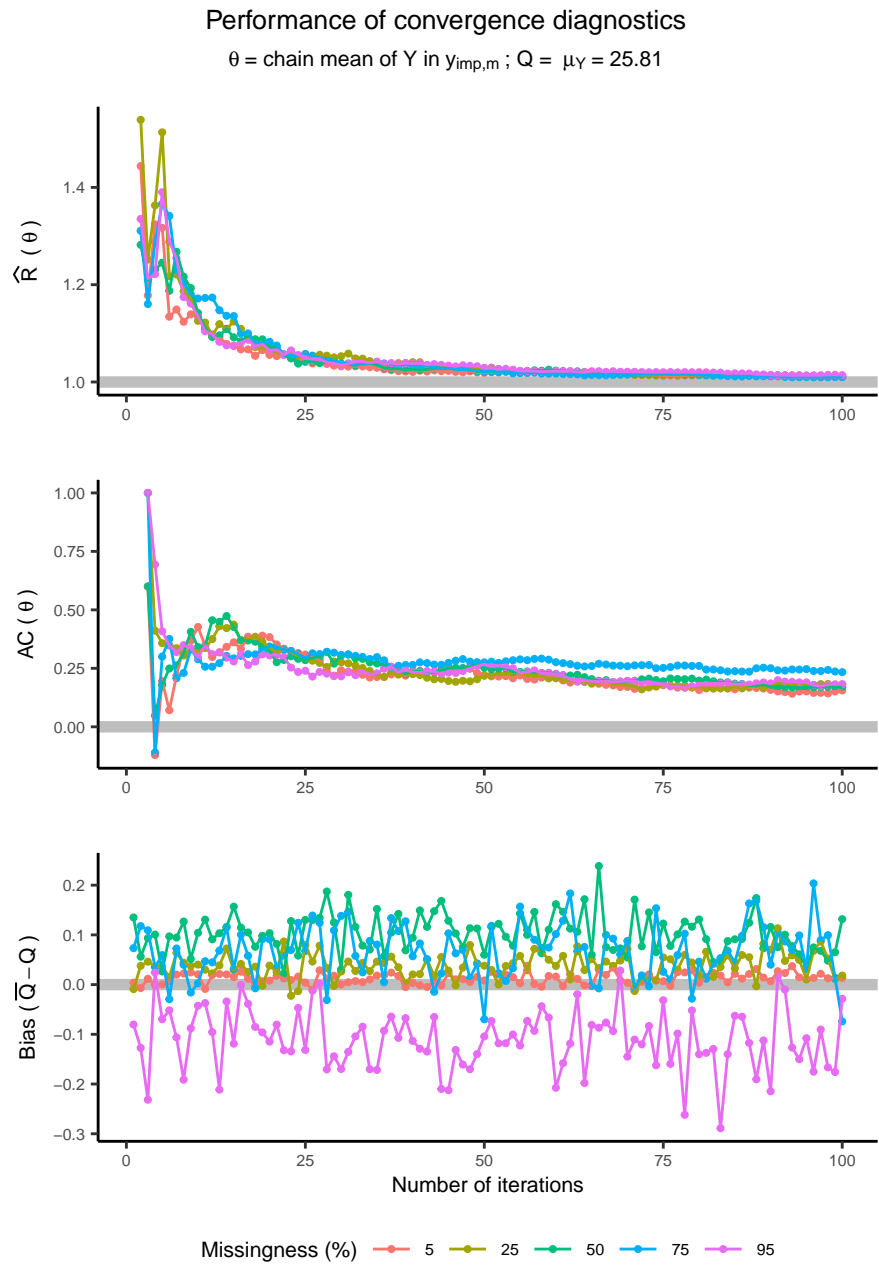


Figure 4. Convergence diagnostics chain mean.

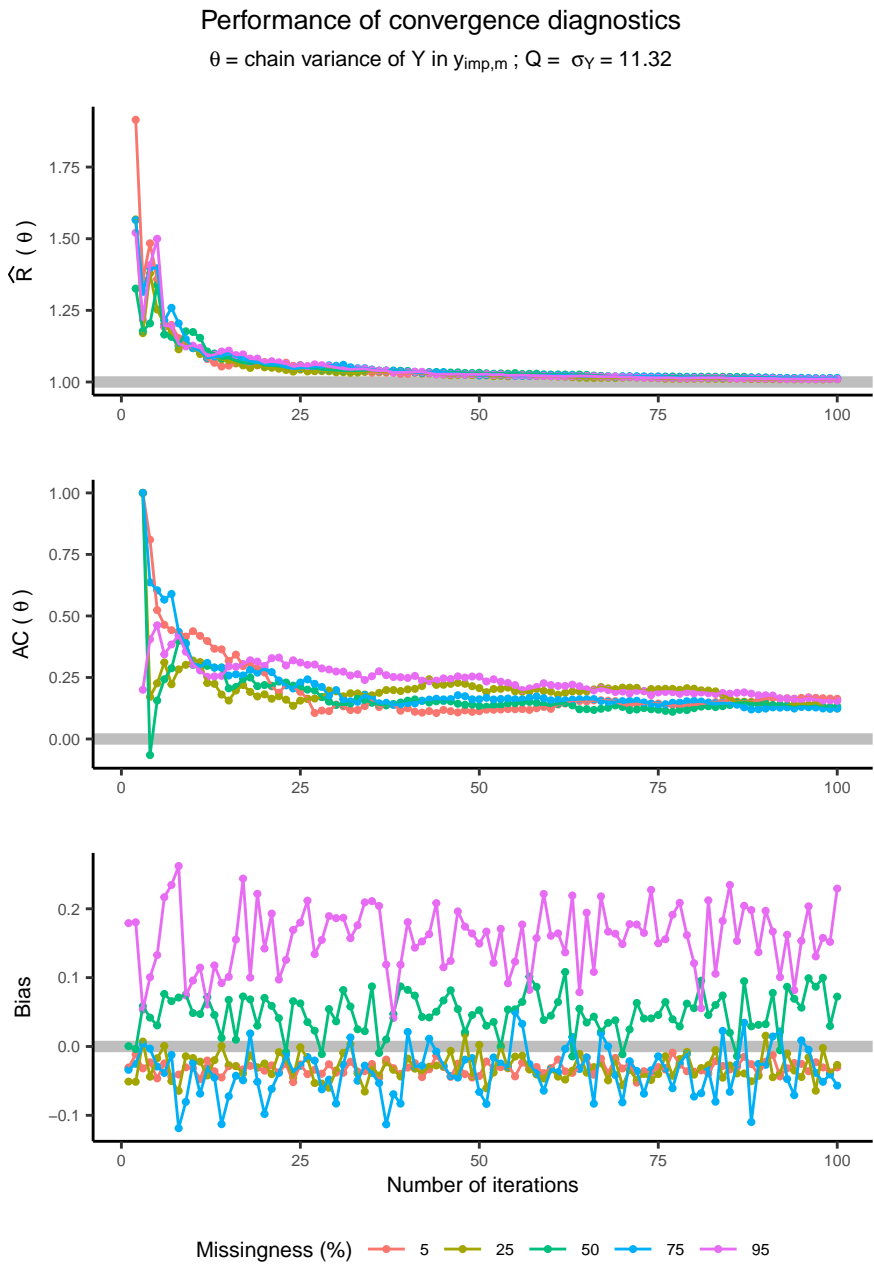


Figure 5. Convergence diagnostics chain variance.




Figure 6. Convergence diagnostics regression coefficient.

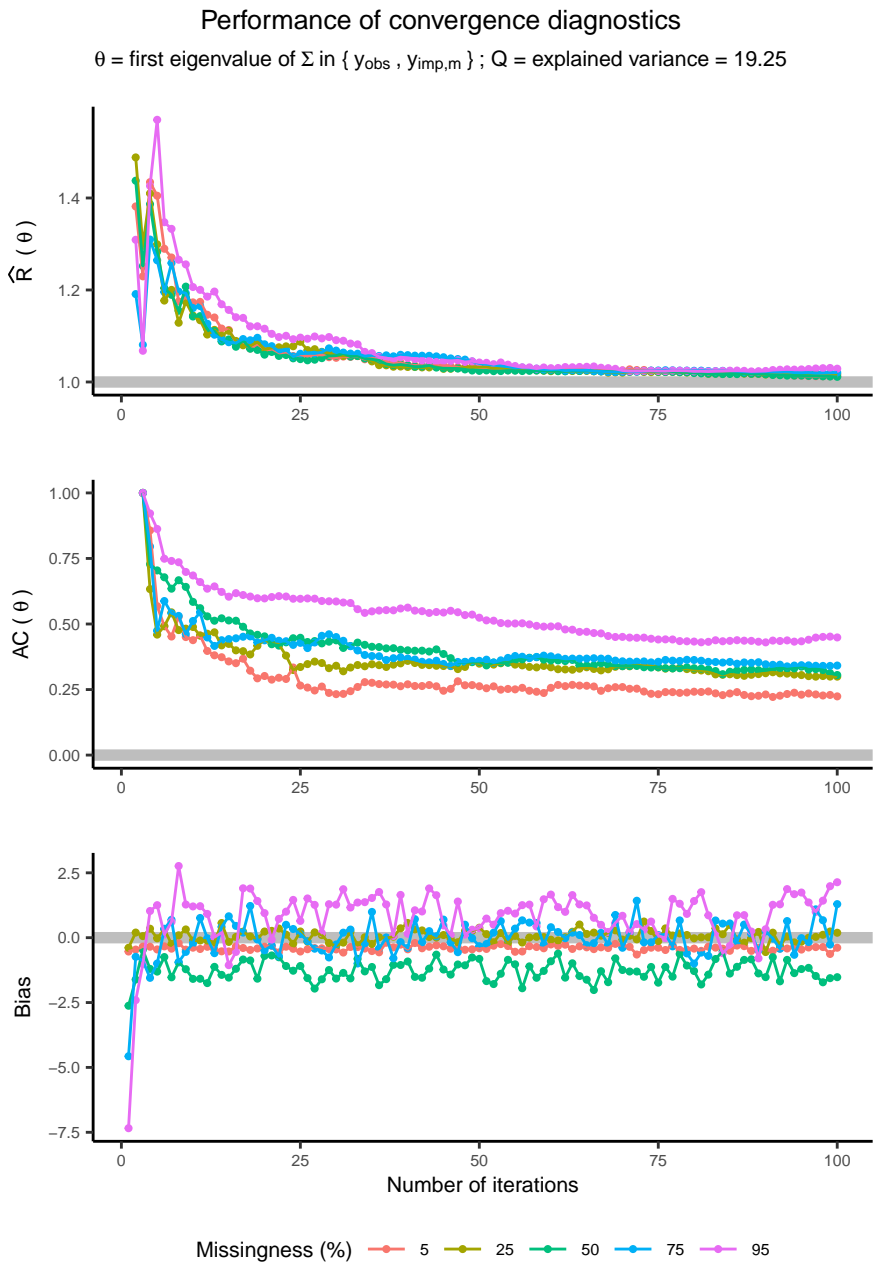


Figure 7. Convergence diagnostics regression coefficient.

conditions, the under-coverage may be attributed to the bias in the estimated regression coefficients.

This is only true for 75% missingness: If we look at the estimated proportion of explained variance in outcome variable Y we see that the coefficient of determination is underestimate estimated in conditions where the number of iterations is equal to two or less, and slightly overestimated in conditions where the number of iterations is equal to three or more.

In short, we see that the minimum number of iterations required to obtain unbiased, confidence valid regression estimates is 5. This value, however, is dependent on the percentage of missing values. E.g., with 95% of cases having missing data we need at least seven iterations to obtain unbiased results.

Discussion

This note shows that convergence diagnostics \hat{R} and AC may diagnose convergence of multiple imputation algorithms, but their performance differs from conventional applications to iterative algorithmic procedures. **(nope! it shows that MICE can lead to correct outcomes when they have not converged according to two common conv diags. This may be due to the measures (e.g., assumption of overdisp) or due to the Qs (lm reg coeff, not higher dimensional/more complex RQs). Add what %miss has to do with it.)**

\hat{R} and autocorrelation indicate that algorithmic convergence may only be reached after twenty or even forty iterations, while unbiased, confidence valid estimates may be obtained with as little as four iterations. These results are in agreement with the simulation hypothesis: \hat{R} over-estimates the severity of non-convergence when applied to MI procedures. %This may be due to the quantity of scientific interest chosen. More ‘complicated’ Qs (e.g., higher-order effects or variance components) might show bias, under- or over-coverage at higher T .

According to this simulation study, the recently proposed threshold of $\hat{R} < 1.01$ may be too stringent for MI algorithms. **(This is only one of the goals: to give applied researchers a diagnostic to indicate that they should keep iterating. The other is the default in mice and other software packages, and yet another is... i forgot)** Under the relatively easy missing data problem of the current study, the threshold was not reached. The other extreme of the \hat{R} -thresholds, the conventionally acceptable $\hat{R} < 1.2$, may be too lenient for MI procedures. Applying this threshold to the current data, lead to falsely diagnosing convergence at $T = 3$ **(because it goes up after, not because it is not converged enough)**. It appears that the widely used threshold of $\hat{R} < 1.1$ suits MI algorithms the best. We might, however, also formulate a new threshold, specifically for the evaluation of MI algorithms. The current study suggests that $\hat{R} < 1.05$ may be implemented, since that is the level at which the \hat{R} stabilize (around $T = 20$) **(Not necessary for this Q, but maybe for more complicated Qs)**.

The negative AC -values obtained in this study show no threat of non-stationarity. However, the initial dip in AC -values **(which disappeared!)** may have implications for the default number of iterations in mice (`maxit` = 5). Terminating the algorithm at

$T = 5$ may not be the most appropriate, since this lead to the worst convergence (**nope, only for Rh, not AC**), as indicated by \hat{R} and AC . Under the current specifications, $T > 20$ would be more appropriate.

Further research is needed to investigate their performance under clear violation of convergence, e.g. dependency between predictors (predictors with very high correlations). Until then, we have only shown that the convergence diagnostics can diagnose non-convergence of MI algorithms that trend towards a converged state. Also for future research, look at developing a convergence diagnostic for substantive models, and implement a Wald test for $AC = 0$.

Recommendations for empirical researchers

For empirical researchers: 1) Check trace plots for pathological non convergence and adjust imputation model if necessary. 2) use \hat{R} wait 1.1 ash threshold and autocorrelation with ? as threshold. Keep iterating until these thresholds are reached. 3) Do not use the R function ACF. Instead, compute autocorrelations manually (see e.g., [GitHub link]). 4) Track your own scalar summary of interest. This is somewhat advanced but explained in Van Buren 2018. Compute \hat{R} and autocorrelation values for this scalar summary. 5) **Something** about the novel θ that is ‘substantive model-independent’.

Recommendations for future research

Pmm, M(N)AR, Empirical data, Vignette, ShinyMICE, AC across imputations, not iterations (new convergence diagnostic unique for iterative imputation??).

If I don’t include it in my study, also: Significance of AC values.

References

- Brooks SP and Gelman A (1998) General Methods for Monitoring Convergence of Iterative Simulations. *Journal of Computational and Graphical Statistics* 7(4): 434–455. DOI: 10.1080/10618600.1998.10474787.
- Cowles MK and Carlin BP (1996) Markov chain Monte Carlo convergence diagnostics: a comparative review. *Journal of the American Statistical Association* 91(434): 883–904.
- Gelman A, Carlin JB, Stern HS, Dunson DB, Vehtari A and Rubin DB (2013) *Bayesian Data Analysis*. Philadelphia, PA, United States: CRC Press LLC.
- Gelman A and Rubin DB (1992) Inference from Iterative Simulation Using Multiple Sequences. *Statistical Science* 7(4): 457–472. DOI:10.1214/ss/1177011136.
- Geweke J (1992) Evaluating the accuracy of sampling-based approaches to the calculations of posterior moments. *Bayesian statistics* 4: 641–649.
- Lacerda M, Ardington C and Leibbrandt M (2007) Sequential regression multiple imputation for incomplete multivariate data using Markov chain Monte Carlo. Technical report, University of Cape Town, South Africa.
- Lynch SM (2007) *Introduction to applied Bayesian statistics and estimation for social scientists*. Springer Science & Business Media.

- MacKay DJ and Mac Kay DJ (2003) *Information theory, inference and learning algorithms*. Cambridge university press.
- Morris TP, White IR and Crowther MJ (2019) Using simulation studies to evaluate statistical methods. *Statistics in Medicine* 38(11): 2074–2102. DOI:10.1002/sim.8086. URL <http://onlinelibrary.wiley.com/doi/10.1002/sim.8086>. Publisher: John Wiley & Sons, Ltd.
- Murray JS (2018) Multiple Imputation: A Review of Practical and Theoretical Findings. *Statistical Science* 33(2): 142–159. DOI:10.1214/18-STS644.
- R Core Team (2020) *R: A language and environment for statistical computing*. Vienna, Austria. URL <https://www.R-project.org/>. Tex.organization: R Foundation for Statistical Computing.
- Raftery AE and Lewis S (1991) How many iterations in the Gibbs sampler? Technical report, Washington University Seattle, Department of Statistics, United States.
- Rubin DB (1976) Inference and Missing Data. *Biometrika* 63(3): 581–592. DOI:10.2307/2335739. URL <https://www.jstor.org/stable/2335739>. Publisher: [Oxford University Press, Biometrika Trust].
- Schafer JL (1997) *Analysis of incomplete multivariate data*. Chapman and Hall/CRC.
- Takahashi M (2017) Statistical Inference in Missing Data by MCMC and Non-MCMC Multiple Imputation Algorithms: Assessing the Effects of Between-Imputation Iterations. *Data Science Journal* 16: 37. DOI:10.5334/dsj-2017-037.
- Van Buuren S (2018) *Flexible imputation of missing data*. Chapman and Hall/CRC.
- Van Buuren S and Groothuis-Oudshoorn K (2011) mice: Multivariate Imputation by Chained Equations in R. *Journal of Statistical Software* 45(1): 1–67. DOI:10.18637/jss.v045.i03. URL <https://www.jstatsoft.org/index.php/jss/article/view/v045i03>.
- Vehtari A, Gelman A, Simpson D, Carpenter B and Bürkner PC (2019) Rank-normalization, folding, and localization: An improved \widehat{R} for assessing convergence of MCMC URL <http://arxiv.org/abs/1903.08008>. ArXiv: 1903.08008.
- Vink G (???) Towards a standardized evaluation of multiple imputation routines .
- Vink G and van Buuren S (2014) Pooling multiple imputations when the sample happens to be the population. *arXiv:1409.8542 [math, stat]* URL <http://arxiv.org/abs/1409.8542>. ArXiv: 1409.8542 version: 1.